

INT3404E 20 - Image Processing: Homework 1

Doan Duc Kien - 21020207

This report presents the implementation of three functions in homework 1, namely `flip_image`, `rotate_image` and `grayscale_image`. Each section includes the corresponding function's code, a brief explanation and an example result.

1 `flip_image`

```
def flip_image(image):  
    """  
    Flip an image horizontally using OpenCV  
    """  
    return cv2.flip(image, 1)
```

The function `flip_image` flips an image horizontally using OpenCV's `cv2.flip` function with a `flipCode` of 1, which means in the horizontal. The result is illustrated in figure 1.



(a) Original image



(b) Flipped image

Figure 1: Result of `flip_image`

2 `rotate_image`

```
def rotate_image(image, angle):  
    """  
    Rotate an image using OpenCV. The angle is in degrees  
    """  
    h, w = image.shape[:2]  
    center = (w / 2, h / 2)  
    M = cv2.getRotationMatrix2D(center, angle, 1.0)  
    rotated = cv2.warpAffine(image, M, (w, h))  
    return rotated
```

To rotate an image by a specific angle, 2 steps are needed:

- 1. Calculate a rotation matrix using `cv2.getRotationMatrix2D`. The matrix is given by:

$$\begin{bmatrix} \alpha & \beta & (1 - \alpha) \cdot \text{center.x} - \beta \cdot \text{center.y} \\ -\beta & \alpha & \beta \cdot \text{center.x} + (1 - \alpha) \cdot \text{center.y} \end{bmatrix}$$

where

$$\alpha = \text{scale} \cdot \cos \text{angle},$$

$$\beta = \text{scale} \cdot \sin \text{angle},$$

$$\text{center} = (x, y) \text{ is the center point of the image}$$

- 2. Apply the transformation given by the rotation matrix using `cv2.warpAffine`.

Figure 2 illustrates the original image after a 45-degree left rotation.



Figure 2: Result of `rotate_image`

3 grayscale_image

```
def grayscale_image(image):
    img_gray = 0.299 * image[:, :, 0:1] + 0.587 * image[:, :, 1:2] + 0.114 * image[:, :, 2:]
    return np.broadcast_to(img_gray.astype(image.dtype), image.shape)
```

To convert an image to grayscale format, we construct an array whose elements are given by a weighted average of RGB values in the original image:

$$p = 0.229R + 0.587G + 0.114B$$

Making use of NumPy's addition operator on arrays, we can directly average the RGB channels. Since the resulting array only has one channel, we broadcast it to make a 3 channel image that has identical values across channels. An example of input and output of this function is shown in figure 3.



(a) Original image



(b) Grayscale image

Figure 3: Result of grayscale_image